



جامعة خليفة
Khalifa University

COSC320 -Assignment 1

Deep Dive into C4 - Understanding a
Self-Interpreting C Compiler

C4 Report

Joud Herzallah

100061269

Due Date: March, 03, 2025



جامعة خليفة
Khalifa University

The lexical analysis process: How does the code identify and tokenize input?

C4's lexer is mostly implemented in the `next()` method. It reads characters from the source code (via a global pointer `p`) and decides the next token based on patterns such as identifiers, numeric literals, string/character constants, operators, and punctuation.

Whitespace and comments (including the `//` and `#` directives) are skipped. When a new identifier is detected, the lexer generates a hash and examines the symbol table to determine if it has previously been saved. Numeric literals can be parsed as decimal, hexadecimal, or octal. Once the token has been discovered, `tk` is set to its type, like (Num, Id, If), and additional fields such as `ival` are filled with numbers or string addresses.

The parsing process: How does the code construct an abstract syntax tree (AST) or equivalent representation?

C4's `expr()` and `stmt()` functions employ a "precedence climbing" or recursive-descent style, in contrast to many other compilers that construct a formal ast. These routines directly output bytecode instructions into an array `e` after parsing statements and expressions, respectively. The parser uses `next()` to read tokens and then creates instructions (such as IMM, ADD, and JMP) that match the parsed constructs depending on the precedence of operators.

By keeping a symbol table (`sym`), local or global variables and function definitions are managed. The `"expr()"` method functions as a mini-ast builder internally and outputs code in postfix form right away, even though the parser in C4 doesn't create a separate ast data structure.

The virtual machine implementation: How does the code execute the compiled instructions?

The final portion of `main()` is where the VM executes. The code sets up registers (`pc`, `sp`, and `bp`) and starts a loop that retrieves and runs each opcode when the parser has filled the instruction array `e`. Arithmetic (ADD, MUL), jumps (JMP, BZ), function calls (JSR, ENT, LEV), loading/storing integers (LI, SI), and system calls (OPEN, READ, etc.) are among the supported instructions.



جامعة خليفة
Khalifa University

In the virtual machine, instructions manipulate the top of a contiguous stack in memory by pushing and popping values. While `sp` (stack pointer) and `bp` (base pointer) monitor local frames, `pc` (program counter) keeps track of which instruction is being performed at any one time.

The memory management approach: How does the code handle memory allocation and deallocation?

`c4` manages memory sparingly. At launch, four big memory pools are set aside for the symbol table (`sym`), code array (`e`), data segment (`data`), and stack (`sp`). These pools are merely used to lay out the compiler's own data structures ex: token storage and symbol table entries. For dynamic allocations at runtime, `c4` provides system calls like `malloc()`, `free()`, and `memset()`, which ultimately defer to the host operating system.

`c4` is a small self-hosted compiler that relies on preallocated pools and the `c` standard library's `malloc/free` for dynamic use. It does not have a comprehensive garbage collector or sophisticated memory scheme.



جامعة خليفة
Khalifa University